

Independent AI evaluations for K-12 EdTech

Fewer than 10% of AI-based approaches in education have undergone independent validation for their efficacy in improving student learning (UNESCO, 2024). OurDojo measures response quality, learning efficacy, and failure-mode transfer on AI-generated educational content.

DIAGNOSTIC *For teams running live AI tooling. Typical engagement: 3 weeks.*

Tailored to your LLM pipeline. We measure your live system against the dimensions that determine perceived quality in K-12 contexts, identify where it's breaking, and leave your team with a sharper internal eval flywheel than they started with. Three deliverables:

- **Error Taxonomy** — failure modes ranked by frequency and severity, mapped against your highest-impact use cases.
- **Evaluation Specifications** — application-specific evals you can run repeatedly, plus a review of your existing metrics and experimentation loop.
- **Prioritized Roadmap** — the changes that buy the most quality per unit of engineering effort.

Worked example. Feb 2026: a teacher-facing content-generation tool, evaluated across 1,000 generated passages using the Learning Commons grade-level appropriateness evaluator. After one diagnosis cycle, the vendor's March re-evaluation under the same protocol showed a **+17pp gain on grade-level accuracy**. Methodology at ourdojo.org/research.

Our commitment — If we cannot deliver clarity on where your system is breaking and a concrete plan to fix it, we continue working at no additional cost until we do.

REPORT *For teams who want an external benchmark on the record. Typical engagement: 2 weeks.*

A 1,000-sample evaluation of your AI output scored against established frameworks from authorities like Student Achievement Partners, CAST, and ANet. Includes an analysis report and 1-hour results walkthrough. Co-publishable when both sides want the result on the record.

LATEST PUBLICATIONS

- *The K-12 AI Quality Playbook*. Apr 2026: A framework for EdTech evaluations with market research and drop-in RFP language for district procurement.
- *Cross-vendor benchmark*. Mar 2026: Twee at 89% grade-level accuracy on Science; a comparable teacher-facing tool at 35% baseline on the same dimension. Standardized evaluation methodology that enables direct comparisons across tools and over time.

Jay Syz, founder. 2+ years building AI evaluation systems at enterprise organizations; previously an engineer at Google. Supported by an expert validation network across staff engineering, data science, and K-12 classroom practice.

syz@ourdojo.org · ourdojo.org